

Homework 5 (due Thursday, March 8th)

By signing my name, I agree to abide by the Stern Code of Conduct _____

Question 1

This question is about the sampling distribution of the mean, and using \bar{X} as an estimator for the parameter μ . Suppose X_1, \dots, X_n are independent and identically distributed (i.i.d.) with the same distribution as X .

a. Recall that $E[\bar{X}] = E[X]$ and $\text{Var}(\bar{X}) = \frac{\text{Var}(X)}{n}$. Do we need to assume anything more about X for \bar{X} to be an unbiased estimate of μ , and if so, what?

b. One way to write Chebyshev's inequality is $P(|Y - E[Y]| \geq k\sqrt{\text{Var}(Y)}) \leq 1/k^2$. Let $Y = \bar{X}$ and write this inequality out in terms of $\bar{X}, E[X], k, \text{Var}(X)$, and n .

c. If we let $k = \sqrt{n}$, the answer from the previous part should simplify to $P(|\bar{X} - E[X]| \geq \sqrt{\text{Var}(X)}) \leq 1/n$. If we want 99% probability that the sample mean is *within one standard deviation* of X away from $E[X]$, how large of a sample do we need, according to this formula?

d. Suppose \bar{X} is unbiased: $E[\bar{X}] = \mu$. What is the mean squared error of this estimate, $\text{MSE}(\bar{X})$? Simplify as much as you can.

e. Assume now that $X \sim \text{Ber}(p)$. For some constant c , $c\bar{X}$ has a distribution that you know. What is the constant, and what is the distribution?

Question 2

This question will guide you through a *simulation study* in R to understand the bias of a certain estimator. Suppose U_1, \dots, U_n are independent and identically distributed as $U[0, \theta]$, with $\theta = 1$ and let $\hat{\theta} = \max\{U_1, \dots, U_n\}$. Since we are generating this data ourselves, we know the true value of $\theta = 1$, so we can compute the bias of $\hat{\theta}$. Our goal will be to study how this bias decreases as the sample size n increases.

The following code creates a *function* that you can use to generate observations of $\hat{\theta}$.

```
theta_hat <- function(n) return(max(runif(n)))
```

You need to run this code once so that R learns the definition of `theta_hat`. After that, you can “call” the function by running, for example, `theta_hat(10)` to generate one observation of a maximum of sample size $n = 10$.

```
theta_hat(10)
```

```
## [1] 0.9485926
```

To generate a sample of many i.i.d. copies of $\hat{\theta}$, we use the `replicate` function:

```
replicate(10, theta_hat(10))
```

```
## [1] 0.8834024 0.9749575 0.7858623 0.8282076 0.9919492 0.9742462 0.7195793
```

```
## [8] 0.8900270 0.8900003 0.7248828
```

Finally, we estimate the bias by generating many samples of $\hat{\theta}$, taking their average, and subtracting θ :

```
mean(replicate(10000, theta_hat(10))) - 1
```

```
## [1] -0.09001053
```

a. Run this code again to estimate the bias when $\hat{\theta}$ is based on a sample of size $n = 100$, and again for a sample of size $n = 1000$.

```
# write code here and delete this comment
```

b. Compute these answers to the answer I gave in class: compute $(n - 1)/n - 1 = -1/n$ for the same values of n . Are the simulation estimates of bias reasonably close to the exact mathematical answer?

c. Use the `sd` function to estimate the standard deviation of $\hat{\theta}$ based on $n = 10$ and on $n = 100$.

```
# write code here and delete this comment
```